

Pricing a Chokepoint: A Synthetic-Control Estimate of the Brent Crude Premium from the 2026 Strait of Hormuz Disruption

Master Project, Data Science for Decision Making
Barcelona School of Economics

June 2026

Abstract

Roughly a fifth of the world's seaborne crude passes through the Strait of Hormuz, so its disruption on 28 February 2026 was widely expected to move global oil prices. Measuring that effect, however, is not straightforward: As Brent crude represents a unique global benchmark, there is no untreated control unit available from which a counterfactual price can be directly inferred. We construct that missing counterfactual with the *synthetic control method* (SCM), building a price path from a pool of 19 macro-financial donor series (metals, agriculturals, equities, foreign exchange, rates, and volatility) chosen to share Brent's exposure to common global factors while staying off the oil-supply causal path. To keep the result from hinging on any single modelling choice, the synthetic is formed as the median of an ensemble of five estimators with different inductive biases. We benchmark the framework using the Russian invasion of Ukraine on 24 February 2022, which represents the closest comparable historical oil supply shock. We then apply the framework to estimate the effects of a disruption in the Strait of Hormuz. Over the post-event window to mid-June 2026 we find an average Brent premium of roughly 56 % (ensemble median; interquartile range 51 % to 57 %), an estimate that is robust to placebo-based inference and a range of robustness and contamination checks. We frame the estimate as a reduced-form, data-driven complement to the structural scenario estimates produced by energy institutions, and discuss its use as an input to strategic-reserve, monetary, and fiscal decisions.

Contents

1	Introduction	3
2	Related Literature	4
2.1	Building counterfactuals for a single aggregate unit	5
2.2	Where SCM has been applied, and the gap we occupy	5
2.3	Oil prices and geopolitical risk: the incumbent methods	5
2.4	Inference with one treated unit and a short post-period	6
3	Conceptual Framework	6
3.1	A factor model of Brent	6
3.2	Potential outcomes and the estimand	7
3.3	Identifying assumptions	7

3.4	The direction of contamination bias	7
3.5	Why an ensemble	8
4	Data	9
4.1	Outcome and sources	9
4.2	The donor pool and the SUTVA audit	9
4.3	Sample windows and summary statistics	10
5	Empirical Strategy and Validation	11
5.1	Estimators	11
5.2	Validation battery and its logic	11
6	Results	12
6.1	Pre-event fit	12
6.2	Headline estimates	12
6.3	Inference: the in-space placebo	13
6.4	Robustness	14
6.5	Interpretation	17
7	Conclusion	17
A	Robustness grid	20
B	Bayesian ridge and full BSTS	20
C	Structural-break tests	20

1 Introduction

The Strait of Hormuz is the narrowest binding constraint in the global oil system. Roughly one-fifth of world petroleum-liquids consumption transits the lane between Iran and Oman, carrying crude from Saudi Arabia, Iraq, the United Arab Emirates, Kuwait, Iran, and Qatar to seaborne markets. What makes the chokepoint so consequential is the absence of a redundant route: the Saudi East–West and UAE Habshan–Fujairah pipelines together divert only a small fraction of this volume overland, and no maritime bypass exists, so the overwhelming majority of Gulf crude has nowhere else to go. This leaves Hormuz the single strongest bottleneck in the world’s oil-transport network. The lane has seen periodic near-misses—the 1980s tanker war, the 2019 Abqaiq strike, recurring tensions through the mid-2020s—each producing a transient risk-premium spike, but never an actual closure. The 28 February 2026 disruption is the first realisation of that long-anticipated tail risk.

Because crude oil is an input to almost every industry and a cost borne across nearly all economies, a disruption on this scale was always going to move prices, and Brent duly rose sharply in the days that followed. We study Brent rather than another benchmark because, while the major crude indices are highly correlated and the choice between them is largely immaterial, Brent is the reference price against which the majority of internationally traded crude is settled; we use its spot quotation rather than a futures contract. Yet the raw price jump is uninformative on its own: it conflates the disruption with everything else moving markets at the same time, and it says nothing about how the effect evolves. This motivates a causal question:

What is the causal effect of the 2026 Strait of Hormuz disruption on the price of Brent crude over a policy-relevant post-event window?

We deliberately target the *total* causal effect rather than any single transmission channel, and we estimate it as a *path* over the post-event window rather than a single number. This choice is driven by the applications: emergency strategic-reserve releases are sized against an implied price displacement; the monetary-policy decision to look through an oil shock or to tighten turns on its size and persistence; fiscal stabilisation in subsidising oil importers scales with the move; and sovereign hedging programmes price disruption scenarios against exactly such magnitudes. Each lever acts on the trajectory of the effect, not on a peak-day figure. The object we need is therefore the average gap between observed Brent and the price that would have prevailed absent the disruption, traced over the whole window. Recovering that counterfactual path, which has no directly observable analogue, is the methodological core of this paper.

Standard observational designs cannot deliver it. Brent is a single, globally traded benchmark, which rules out difference-in-differences (no panel of comparable treated and control units), regression discontinuity (no running variable around a cutoff), and instrumental variables (no credible instrument for a one-off geopolitical event). The *synthetic control method* (Abadie et al., 2010; Abadie, 2021) is built for exactly this situation: it forms a weighted basket of donor series whose pre-event co-movement with the outcome is high and projects that basket forward as the counterfactual. It yields a full counterfactual path rather than a point estimate, makes its identifying assumptions auditable through the donor pool, and requires no structural stance on whether the move is supply, risk premium, or expectations—its estimand is a reduced-form total effect.

The tools most commonly used to analyze such price shocks include event studies, structural VARs, and institutional scenario models. However, they tend to address narrower or different questions. Event studies measure abnormal returns in a short window around the event but yield no multi-month path and miss persistence. Structural VARs (Kilian, 2009; Baumeister and Hamilton, 2019) decompose prices into supply, demand, and inventory shocks but produce no counterfactual path and rely on identification restrictions with no settled precedent for a

chokepoint event. And institutional scenario models propagate an assumed barrel shortfall through a supply–demand model, so their counterfactual is imposed rather than read from observed prices. We see these as complements rather than competitors: our cross-asset, data-driven estimate and a structural scenario estimate answer different questions, the price gap implied by observed co-movement versus the price implied by an assumed shortfall, so their agreement is itself informative, and in practice the two can be used side by side.

Methodologically, our closest neighbour is the SCM-on-price design of [Becker et al. \(2021\)](#). That work studies a regulated retail-fuel market with cross-country donor units; neither is available here, since Brent is a global crude benchmark and no untreated country or market can serve as a donor. We therefore replace geographic donors with cross-asset donors, including metals, agriculturals, equities, foreign exchange, rates, and volatility, and apply the design under an exogenous geopolitical treatment. To our knowledge this combination has not been used for any previous oil shock, so the contribution is as much about establishing whether such a design is valid in this setting as about the number it produces.

A single estimate, however, cannot be told apart from an artefact of the donor pool, pre-window, model class, and regularisation. We therefore benchmark the design on the nearest comparable historical oil-supply shock, the 2022 Russian invasion of Ukraine (Brent \$95→\$130 within two weeks), and check that the same machinery recovers that documented magnitude before estimating the Hormuz effect afresh on its own pre-window. We do not transfer the Russia fit to Hormuz: the two events sit in very different market regimes and each is estimated independently. Because both are fit on a common 19-series donor pool, we can additionally compare the donor weights across regimes as a check on the stability of the factor structure SCM relies on.

To prevent any single modelling choice from driving the result, the synthetic is formed as the median of an ensemble of five estimators with different inductive biases: convex SCM, augmented SCM, elastic net, gradient-boosted trees, and Bayesian ridge. Fitting a synthetic that tracks pre-event Brent closely is the easy part; the substantive work of this paper is the validation that challenges the donor pool and the no-interference assumption it rests on—placebo-based inference, leave-one-donor-out checks, structural-break tests, a signed contamination-bias analysis, and a comparison against naïve univariate counterfactuals. Our contribution is thus twofold: a data-driven Brent-price ATT for the Hormuz disruption that complements existing structural estimates, and a transparent template for validating cross-asset synthetic control under a single-unit geopolitical treatment.

The remainder of the paper reviews the relevant literature (Section 2), sets out the factor model, estimand, and identifying assumptions (Section 3), describes the data and donor audit (Section 4), details the empirical strategy and validation battery (Section 5), and reports results for both events (Section 6) before concluding (Section 7).

2 Related Literature

We organise the review around problems rather than individual papers, because our contribution sits at the intersection of three otherwise separate strands: how causal counterfactuals are built for a single aggregate unit, how the oil-price effect of geopolitical shocks is conventionally measured, and how inference is conducted with one treated unit and a short post-period. Tools borrowed from machine learning and time series are introduced where they are used (Section 5 and the appendices); here we position the causal design.

2.1 Building counterfactuals for a single aggregate unit

The problem of constructing a counterfactual for one treated unit is what synthetic control was invented to solve. The method originates with [Abadie and Gardeazabal \(2003\)](#) and [Abadie et al. \(2010\)](#), who build the counterfactual as a convex combination of untreated donors matched on pre-treatment outcomes and formalise permutation (placebo) inference for the single-unit case; [Abadie et al. \(2015\)](#) extend it to comparative politics, and [Abadie \(2021\)](#) codifies the practical requirements we adopt: a long, shock-free pre-window, donors clean of the treatment (the SCM analogue of SUTVA), and a good pre-treatment fit as a precondition for credibility. A subsequent literature relaxes the convex-hull restriction: the augmented SCM of [Ben-Michael et al. \(2021\)](#) adds a ridge bias-correction permitting limited extrapolation when the treated unit lies outside the donor hull, [Doudchenko and Imbens \(2016\)](#) embed SCM, difference-in-differences, and regression in a common regularised-regression family that allows unconstrained weights, and [Athey et al. \(2021\)](#) recast the problem as matrix completion. A caution runs through all of it: convex SCM weights are *non-unique* when donors are correlated ([Abadie and L’Hour, 2021](#)), so many weight vectors yield near-identical fit and individual weights should not be over-interpreted. We use this entire toolbox by running an ensemble that spans the convex, augmented, and regularised-regression cases, and we read cross-model weight dispersion through the non-uniqueness result rather than as instability.

2.2 Where SCM has been applied, and the gap we occupy

Despite this methodological maturity, SCM has rarely been applied to a daily market price. Most peer-reviewed applications estimate effects on *non-price* aggregates such as GDP, trade, and employment, including under discrete geopolitical shocks: [Gharehgozli \(2017\)](#) estimates the GDP cost of sanctions on Iran and [Mancini et al. \(2024\)](#) the trade effect of the 2022 Russia–Ukraine war. The closest mechanical precedent to a price application is [Becker et al. \(2021\)](#), who apply SCM to Austrian retail fuel prices with other European countries as donors; that design differs from ours in three ways that matter, as its treatment is an endogenous regulation rather than an exogenous shock, its outcome an administered retail price rather than a traded benchmark, and its donors cross-country rather than cross-asset. [Juárez et al. \(2024\)](#) use an oil-price shock as the treatment but on regional GDP, not on a price. The intersection we occupy—SCM on a crude-benchmark *price* under an exogenous *geopolitical* disruption, with cross-asset donors—has no direct precedent, and that is the gap this thesis addresses. Because the Hormuz event is recent, the only existing analysis is institutional ([Kiel Institute for the World Economy, 2026](#); [Oxford Institute for Energy Studies, 2026](#); [World Bank, 2026](#); [International Energy Agency, 2026](#); [Federal Reserve Bank of Dallas, 2026](#)); these share one structure—assume a barrels-per-day shortfall, propagate it through a structural model, and output a price path—and none applies synthetic control to the observed price. They form the comparison set against which we position, and complement, our estimate.

2.3 Oil prices and geopolitical risk: the incumbent methods

The methods conventionally used for oil-price effects answer narrower or different questions than ours. The dominant approach is the structural VAR with identified supply, demand, and inventory shocks ([Kilian, 2009](#); [Baumeister and Hamilton, 2019](#)), with [Hamilton \(2009\)](#) the narrative anchor for 2007–08; these decompose price changes into structural components but yield no counterfactual path and require identification restrictions for which a chokepoint event offers no standard precedent. A natural alternative would regress oil prices on the Geopolitical Risk index of [Caldara and Iacoviello \(2022\)](#), but this recovers the *average* response to a unit of geopolitical risk pooled across many events—a different object from the effect of one specific

disruption—and the index spikes endogenously to the very event we study, so it can serve neither as our estimand nor as a clean donor. Short-window event studies (Brown and Warner, 1985) complete the incumbent set and give a magnitude but not the multi-month path policy requires.

2.4 Inference with one treated unit and a short post-period

Our setting makes inference non-standard, and the recent literature counsels restraint in two places. First, on inference itself: permutation/placebo tests (Abadie et al., 2010, 2015) are the field standard, with validity resting on an approximate-symmetry assumption (Canay et al., 2017); Hahn and Shi (2017) show that for convex SCM this effectively requires Gaussian errors, which bounds how far placebo p -values can be pushed for our nonlinear ensemble members, and Chen and Yan (2023) formalise the in-time placebo into a mixed test with a permutation p -value. Second, on pretesting: Roth (2022) shows that pre-trend tests have low power at applied sample sizes and can manufacture false confidence, and Rambachan and Roth (2023) propose bounding rather than testing such violations; we accordingly adopt a report-and-interpret stance toward pre-period diagnostics rather than imposing pass/fail thresholds. Finally, the value of a donor depends on how reliably it tracks the common factor: a literature on matching reliability before differencing shows that a good pre-period fit can *mask* rather than remove contamination, the logic we use to split donors by pre-window factor R^2 when signing the contamination bias (Section 6). Multiple testing across donors is controlled by Benjamini–Hochberg (Benjamini and Hochberg, 1995).

3 Conceptual Framework

This section sets out the data-generating process we assume for Brent, derives the estimand and identifying assumptions from it, and explains why an *ensemble* of estimators is the appropriate proxy when the functional form of the counterfactual is unknown. The factor model in Section 3.1 is the organising device: it makes donor selection an economic exercise, renders each identifying assumption transparent, and gives the contamination bias an explicit sign.

3.1 A factor model of Brent

We decompose log Brent into a component spanned by global macro-financial factors common to many assets and an oil-specific component:

$$p_t^{\text{Brent}} = \underbrace{\lambda' F_t}_{\text{common global factors}} + \underbrace{o_t}_{\text{oil-specific (supply, chokepoint, geopolitics)}} + \varepsilon_t. \quad (1)$$

F_t collects the five drivers that move Brent in normal times—the global industrial cycle, the US dollar, real interest rates, global risk appetite, and China demand—each proxied by non-oil assets. The oil-specific term o_t is the object of interest: the supply and chokepoint component a Hormuz disruption shifts. Each donor j admits its own factor representation $X_{jt} = \gamma_j' F_t + u_{jt}$ with, by construction, *zero loading on* o_t : a clean donor is exposed to the same global factors as Brent but is not physically routed through oil supply. A synthetic control is then a weighted donor basket $\hat{p}_t = \sum_j w_j X_{jt}$, and the logic of the method follows directly from (1). If the weights reproduce Brent’s factor loadings, $\sum_j w_j \gamma_j \approx \lambda$, the basket tracks Brent through the pre-window, where o_t is small and stable. When a chokepoint shock fires, o_t jumps for Brent but not for the basket; the two decouple, and that decoupling *is* the estimated premium. Donor selection is therefore the economic exercise of spanning F_t while excluding o_t —choosing identification over fit. This is precisely why the energy complex (WTI, refined products, petrocurrencies), despite

the highest raw correlation with Brent, is excluded by design: it loads on the very o_t we are trying to measure.

3.2 Potential outcomes and the estimand

Equation (1) describes the observed price; the causal question concerns the price that would have prevailed without the event. Let Y_t^N be Brent’s potential log-price absent the disruption and Y_t^I under it, with treatment beginning at T_0 . The treated-unit effect is $\tau_t = Y_t^I - Y_t^N$ for $t \geq T_0$. With only one treated unit, Y_t^N is never observed after T_0 ; the synthetic supplies the missing term, $\hat{Y}_t^N = \sum_j w_j X_{jt}$, and the estimated effect is the gap $\hat{\tau}_t = Y_t^I - \hat{Y}_t^N$, reported in percent as $\text{gap}_t = 100 [\exp(Y_t^I - \hat{Y}_t^N) - 1]$. We summarise the post-window by the *average treatment effect on the treated*,

$$\widehat{\text{ATT}} = \frac{1}{|W|} \sum_{t \in [T_0, T_0+W]} \text{gap}_t. \quad (2)$$

We prefer the ATT to the peak or terminal gap because it integrates over the entire path, including any reversion—the object that reserve sizing, monetary pass-through, and fiscal buffering act on—and because averaging dampens single-day jumps from liquidity or futures-roll mechanics.

3.3 Identifying assumptions

Because the synthetic stands in for an unobserved counterfactual, the estimate is only as credible as the assumptions that license it, and the factor model makes each one explicit and checkable.

1. **Common-factor spanning (pre-period fit).** The donor basket must reproduce Brent’s factor loadings, so that good pre-event tracking reflects shared F_t rather than coincidence. Checked by pre-period RMSPE, walk-forward cross-validation, and moment matching.
2. **Donor cleanliness (SUTVA / no interference).** No donor may carry the treatment: γ_j must not load on o_t , and donor j must not be hit through a channel specific to the focal event. Defended by an a-priori audit and confirmed by model-agnostic break tests at T_0 .
3. **Regime stability.** The loadings λ, γ_j must hold across the pre-window and the projection window, so that a basket which tracked Brent before T_0 remains a valid counterfactual after it. Checked by distance-correlation and structural-break diagnostics, and, across events, by comparing donor weights between regimes.

Assumptions 1 and 3 are testable from observed data; assumption 2 is the genuinely untestable one, since donor cleanliness concerns the unobserved X_{jt}^N . The next subsection shows how the factor model nonetheless lets us sign the *direction* in which a violation of assumption 2 would move the estimate.

3.4 The direction of contamination bias

The factor representation turns the informal worry that “the donors might be affected by the war” into a quantity with a definite sign. Writing $\delta_{jt} = X_{jt} - X_{jt}^N$ for the event-induced component of donor j , the estimation error is, as an exact identity,

$$\hat{\tau}_t - \tau_t = Y_t^I - \hat{Y}_t^N = - \sum_j w_j \delta_{jt}. \quad (3)$$

Table 1: The five-estimator ensemble. Each contributes a distinct inductive bias; the headline is the cross-model median gap.

Estimator	Inductive bias	Reference
Convex SCM	Non-negative donor weights summing to one; synthetic inside the donor convex hull	Abadie et al. (2010)
Augmented SCM	Convex SCM plus ridge bias-correction; limited extrapolation outside the hull	Ben-Michael et al. (2021)
Elastic net	Signed, sparse+ridge regression; lasso term drops donors entirely	Doudchenko and Imbens (2016)
Gradient boosting (XGBoost)	Nonlinear, tree-based; captures interactions; cannot extrapolate beyond the training range	Chen and Guestrin (2016)
Bayesian ridge	Bayesian linear regression on donor levels (regression component of BSTS); conjugate Gaussian-Gamma prior	Brodersen et al. (2015)

The bias is *minus* the weight-averaged donor contamination. Its sign is therefore an empirical matter, not an assumption: if contaminated donors are pushed *up* (oil-linked terms-of-trade, inflation→rates), the synthetic is pulled up and the ATT is *under*-estimated (conservative); if pushed *down* (demand crowd-out, risk-off), the ATT is over-estimated. The identity (3) is exact, but δ_{jt} is itself unobserved, so signing the bias in practice requires estimating each δ_{jt} against a Brent-free reference—a deliberately naïve construction whose assumptions and limits we set out, and apply, in Section 6. What the theory delivers here is the *structure*: which way a given contamination pushes the estimate, not a certified magnitude. The same identity also distinguishes two regimes of contamination by timing. *Contemporaneous* contamination occurs at T_0 and is screened by the audit and break tests; *delayed* contamination accrues over the post-window, enters through the projection, biases the gap toward zero, and grows with horizon—so a gap that *declines* as the window lengthens is its signature, and the short-horizon gap is the least-contaminated estimate.

3.5 Why an ensemble

A single specification cannot be trusted because the counterfactual’s functional form is unknown: equation (1) need not be linear in the donors, and the convex-hull restriction of classical SCM may or may not bind. Rather than commit to one form, we proxy the counterfactual with five estimators whose inductive biases bracket the plausible cases (Table 1): a simplex-constrained average, a ridge-extrapolating correction, a sparse signed regression, a nonlinear tree ensemble, and a Bayesian linear model. An estimated premium that survives across biases this different is far more likely to reflect a real effect than a single-model artefact. We report the cross-model *median* as the headline—robust to any one model misbehaving—and the inter-quartile range as a model-uncertainty band that complements each model’s own within-model interval.

A note on labelling: the fifth estimator is a Bayesian *ridge* regression, not full Bayesian structural time series. It implements only the regression-on-donors component $\beta'x_t$ of Brodersen et al. (2015) and omits the trend, seasonal, and spike-and-slab selection components; at our sample size, and with no strong weekly seasonality, the regression term carries most of the signal, but the model is best read as a Bayesian-flavoured diversifier rather than a state-space causal-impact estimator (Appendix B).

4 Data

This section describes the outcome series, the donor candidates and the per-event cleanliness audit that selects among them, and the sample windows over which each event is fit and projected. Two themes recur and both follow from the factor model of Section 3: donor selection is an identification choice rather than a fit-maximising one, and the length of the estimation window is itself a SUTVA consideration, not merely a sample-size one.

4.1 Outcome and sources

The outcome is the daily *Europe Brent Spot Price FOB* (EIA series RBRTed), in USD per barrel, downloaded directly from the U.S. Energy Information Administration ([U.S. Energy Information Administration, 2026](#)). The donor candidates are 32 macro-financial series from Yahoo Finance (futures, ETFs, FX, rates, volatility), updated incrementally so the panel keeps pace with new data. Table 2 summarises the sources; all are public and redistributable, and the panel is rebuilt from source on each run so the analysis is reproducible against the latest available data (here, through 15 June 2026 for Brent, the binding constraint given the EIA’s roughly one-week publication lag, and 23 June 2026 for the donors).

Table 2: Data sources. All series are daily and publicly redistributable.

Object	Source	Contents
Brent spot (outcome)	EIA, series RBRTed	Europe Brent Spot FOB, USD/bbl, 1987–present
Donor candidates (32)	Yahoo Finance	Metals, agriculturals, equities, FX, rates, credit, volatility, crypto

4.2 The donor pool and the SUTVA audit

Donor selection is central to identification (Section 3): a series qualifies only if it (i) co-moves with Brent through common global factors, (ii) is *not* physically routed through the disruption being attributed, (iii) is liquid and daily-traded, and (iv) has sufficient pre-history. Criterion (ii) is the SUTVA (no-interference) condition and is enforced by a per-event audit that classifies each donor as Clean, Mild, or Heavy. For **Russia 2022**, 13 of 32 donors are flagged through identifiable channels: Russia/Ukraine grain exports (Wheat, Corn), palladium supply concentration, the EU energy-crisis pass-through to EUR and the dollar basket, and a sanctions-evasion narrative for Bitcoin. For **Hormuz 2026**, only Cotton is flagged (the indirect oil→polyester substitution channel); no donor is physically routed through the Strait.

To make the two events directly comparable, we fit both on a single shared 19-donor pool, formed by retaining only the donors rated Clean for Russia—the more demanding of the two audits—which reduces the 32 candidates to 19 and, because the Russia strict-clean set is itself a subset of the Hormuz clean set, leaves every member clean for both events. Cotton, flagged Mild for Hormuz, is therefore excluded on the strength of the Russia audit rather than retained, so its indirect channel never enters the shared pool. One donor, VIX, is borderline: a full-history Bai–Perron test ([Bai and Perron, 1998](#)) flags a single break in its relationship with Brent on 2022-01-07 (~7 weeks before the invasion). We treat that as a loading-stability signal rather than a SUTVA violation, because it does not reproduce when the test is run on the analysis window alone, and a sensitivity check shows that keeping or dropping VIX moves the ensemble premium by less than 1 percentage point. We therefore retain VIX, which keeps the volatility factor represented. The energy complex (WTI, refined products, petrocurrencies) is excluded by

construction: it has the highest raw correlation with Brent but loads on the very oil-specific shock we measure. Table 3 lists the 19-donor pool by factor category.

Table 3: The shared 19-donor pool, by factor category. The full 32-donor pool adds Russia-contaminated series (e.g. Wheat, Corn, Palladium, EUR, DXY, BTC, Copper, Iron Ore, US10Y) used only in the appendix robustness pool and as a contamination diagnostic.

Category	Donors (shared pool)
Metals (3)	Silver, Platinum, Gold
Agriculturals (3)	Coffee, Sugar, Live Cattle
Equities (2)	S&P 500, Nikkei 225
Foreign exchange (8)	AUD, JPY, CHF, CNY, INR, KRW, ZAR, MXN
Rates & credit (2)	TLT (long Treasuries), HYG (high-yield credit)
Volatility (1)	VIX (implied S&P 500 volatility)

4.3 Sample windows and summary statistics

The estimation window must be chosen, not maximised. For daily-traded benchmarks like Brent and its donors, a longer pre-history is not automatically better: these series pass through distinct market regimes, and a window that spans several of them risks teaching the synthetic a donor–Brent relationship that no longer holds in the projection period—a regime shift that re-enters as post-event interference and thus as a SUTVA violation. Our first task was therefore to locate the longest pre-window that contains *no* prior shock or regime break, rather than the longest pre-window available. Each event then has such a pre-window (for fitting) and a post-window (for projecting the counterfactual), each ~ 20 months. Following [Abadie \(2021\)](#), the Russia pre-window starts 2020-07-01, deliberately *after* the March–May 2020 cluster (OPEC+ collapse, WTI-negative, acute COVID) that would otherwise impose an abnormal factor loading. The Russia post-window is truncated at 2022-09-30, the last business day before OPEC+’s October 2022 supply cut—a Brent-specific policy shock that would confound the invasion premium. The Hormuz post-window extends to the latest available data (2026-06-15). Table 4 reports summary statistics; Figure 1 plots the full Brent series with both events.

Table 4: Summary statistics of Brent spot (USD/barrel) over each event sample. Pre/post split at T_0 . Russia: $T_0=2022-02-24$; Hormuz: $T_0=2026-02-28$.

Event	Segment	n (days)	Mean	SD	Min	Max
Russia 2022	Full sample	572	75.95	24.76	36.33	133.18
	Pre-window	421	64.26	16.39	36.33	101.66
	Post-window	151	108.54	11.37	82.55	133.18
Hormuz 2026	Full sample	515	76.95	14.30	59.93	138.21
	Pre-window	443	72.08	6.58	59.93	88.66
	Post-window	72	106.92	12.31	77.24	138.21

The two pre-windows differ in a way that matters later. The Russia pre-period is volatile (SD \$16.4, spanning the COVID recovery and the 2021 reflation rally); the Hormuz pre-period is much calmer (SD \$6.6). This difference drives much of the contrast in the two events’ inference: a volatile pre-period widens the placebo distribution and weakens Russia’s permutation test, even though its magnitude is well-documented.




Figure 1: Brent crude spot price, 2020–2026, with the two focal event dates. Brent roughly doubled after both shocks from a similar $\sim \$75$ base, but from very different pre-event trajectories: a steep 2020–22 recovery rally before Russia, and a mild net decline through 2024–25 before Hormuz. This trajectory asymmetry is what makes a naïve trend-extrapolation counterfactual unreliable (Section 6).

5 Empirical Strategy and Validation

5.1 Estimators

All five estimators (Table 1) take the same input (log donor levels) and target log Brent over the pre-window, then project the post-window counterfactual. The two events are estimated entirely separately: they share the same 19-donor pool and the same five-estimator pipeline, but each event trains its own models on its own pre-window, so no parameter learned on one event is carried to the other. *Convex SCM* solves the classic nested optimisation: non-negative weights summing to one, chosen to minimise pre-period RMSPE (Abadie et al., 2010). *Augmented SCM* adds a ridge penalty that corrects for imperfect pre-fit and allows limited extrapolation (Ben-Michael et al., 2021). The *elastic net* relaxes the simplex to signed, sparse coefficients (Zou and Hastie, 2005; Doudchenko and Imbens, 2016). *XGBoost* fits shallow gradient-boosted trees (Chen and Guestrin, 2016), deliberately constrained (depth-2 stumps, strong regularisation, row/column subsampling) for the small-sample regime. *Bayesian ridge* is a conjugate Bayesian linear regression on donor levels, the regression component of BSTS (Brodersen et al., 2015). Hyperparameters are tuned over small grids (3–9 configurations) to bound the selection bias of Cawley and Talbot (2010); the convex SCM has no tuned hyperparameter. Because the five estimators produce incommensurable weight vectors—a simplex for convex SCM, signed unconstrained coefficients for the regressions, and no linear weights at all for XGBoost—the ensemble operates on the *gap series* each model implies rather than on the weights: the headline at every t is the cross-model median gap, with the inter-quartile range as the model-uncertainty band.

5.2 Validation battery and its logic

A credible SCM estimate must clear three levels (Abadie, 2021): pre-event fit, donor identification, and inference. We organise the validation around these three levels and run every test per model and per event.

The first level concerns *fit*. Out-of-sample tracking is assessed by expanding-window walk-forward cross-validation (Hyndman and Athanasopoulos, 2018; Bergmeir and Benítez, 2012). Its role here is not model selection but evidence against over-fitting: a synthetic whose out-of-sample pre-period tracking matches its in-sample tracking is reproducing a stable donor–Brent relationship rather than memorising the pre-window, which is the empirical content of the common-factor spanning assumption.

The second level concerns *identification*, and is the donor-cleanliness battery that defends the no-interference (SUTVA) assumption. Before any model is fit, each donor’s own return series is tested for a mean shift at T_0 using a permutation form of the known-break test of Chow (1960) (rather than the unknown-break test of Andrews (1993), since the date is known), together with a Wilcoxon event-window test (Brown and Warner, 1985); a donor is flagged if either rejects after Benjamini–Hochberg FDR control (Benjamini and Hochberg, 1995). Regime stability is checked with distance correlation across pre-period thirds (Székely et al., 2007) and with Bai–Perron break dating (Bai and Perron, 1998, 2003), and an I(0)/I(1)-robust trend-break test (Harvey et al., 2009; Kwiatkowski et al., 1992) confirms that no donor’s own trend breaks at either T_0 ; the single VIX relationship-break noted in Section 4 falls ~ 7 weeks *before* the Russia T_0 , not at it, and does not recur on the analysis window. Empirically, for Hormuz 31 of 32 donors agree between audit and tests (only Cotton differs); for Russia the tests confirm the broad picture and add one flag (CNY) attributable to concurrent China dynamics rather than the war.

The third level concerns *inference*. We run the in-space placebo, refitting each donor as a placebo treated unit and ranking Brent’s post/pre RMSPE ratio against the resulting distribution (Abadie et al., 2010), reading the permutation p -value with the symmetry caveat of Hahn and Shi (2017); the in-time placebo at a fake T_0 six months early, formalised as the mixed placebo test of Chen and Yan (2023); and leave-one-donor-out. Two further benchmarks sit alongside the formal tests: a deliberately contaminated set of naïve univariate counterfactuals (random walk, drift, linear trend, and an AIC-selected ARIMA that reduces to a driftless random walk in both events) as a floor, and the signed contamination-bias diagnostic of equation (3). Following Roth (2022); Rambachan and Roth (2023), pre-period diagnostics are reported and interpreted rather than turned into pass/fail thresholds; flagged models are *not* auto-excluded, because the cross-model IQR is itself the model-uncertainty band.

6 Results

6.1 Pre-event fit

Table 5 reports walk-forward CV. For Hormuz, four of five models clear the heuristic flag (validation/training RMSE ratio below 2); only Bayesian ridge, at 2.28, is flagged. For Russia, two of five clear it (convex SCM and elastic net); with VIX kept in the pool the augmented SCM rises just past the threshold (2.06), and XGBoost (3.24) and Bayesian ridge (3.09) fail more clearly. The flags are not exclusions: they identify which models to read with caution, and the most flagged model (Bayesian ridge) is the one whose headline gap is the ensemble outlier.

6.2 Headline estimates

Table 6 reports the per-model post-window gap, the drift contribution, and the drift-adjusted gap, for both events; Figure 2 plots the observed and synthetic paths. **Russia 2022** yields an ensemble-median premium of 22.0% over the seven-month window (IQR 10.9% to 30.8%), with an implied counterfactual mean of \$88.6/bbl against an actual \$108.5. This is consistent with the documented \$95→\$130 move and its partial reversion by September. **Hormuz 2026** yields an ensemble-median premium of 55.9% over the window to mid-June (IQR 51.2% to

Table 5: Walk-forward cross-validation (5 expanding folds, 20-day horizons). A validation/training RMSE ratio above 2 is a heuristic overfit flag, not an exclusion rule.

Event	Model	train RMSE	val RMSE	val/train	Flag
Russia	Convex SCM	0.113	0.141	1.25	OK
Russia	Augmented SCM	0.046	0.095	2.06	overfit
Russia	Elastic net	0.051	0.102	1.99	OK
Russia	XGBoost	0.039	0.127	3.24	overfit
Russia	Bayesian ridge	0.034	0.105	3.09	overfit
Hormuz	Convex SCM	0.060	0.083	1.40	OK
Hormuz	Augmented SCM	0.049	0.075	1.54	OK
Hormuz	Elastic net	0.052	0.079	1.51	OK
Hormuz	XGBoost	0.044	0.082	1.84	OK
Hormuz	Bayesian ridge	0.034	0.078	2.28	overfit

56.6 %), with an actual post-window mean of \$106.9 against an implied counterfactual of $\sim$ \$68.6. The Hormuz IQR is far tighter than Russia’s ($\sim$ 5 vs $\sim$ 20 pp), reflecting the calmer Hormuz pre-period.

Table 6: Headline post-event gaps (%). Drift-adjusted = raw – drift contribution. Ensemble = cross-model median; band = IQR (Q25–Q75).

Model	Russia 2022 (7 mo)			Hormuz 2026 ($\sim$ 3.5 mo)		
	Raw	Drift	Adj.	Raw	Drift	Adj.
Convex SCM	33.6	+5.2	28.4	55.9	−1.4	57.3
Augmented SCM	10.9	+0.3	10.6	56.8	−0.4	57.1
Elastic net	22.0	+0.7	21.2	56.6	−0.8	57.4
XGBoost	30.8	+1.8	28.9	51.2	−1.5	52.7
Bayesian ridge	−9.0	+0.1	−9.1	43.0	−0.0	43.0
Ensemble median	22.0	n/a	21.2	55.9	n/a	57.1
IQR	[10.9, 30.8]	n/a	n/a	[51.2, 56.6]	n/a	n/a

Economic reading of the model spread. The cross-model dispersion has an economic explanation rather than being noise. Convex SCM, confined to the simplex, leans on the most stable donors (the FX block) as a “trend baseline” and adds soft-commodity deltas (Coffee, Sugar for Russia). The Bayesian ridge, with unconstrained Gaussian coefficients, loads heavily and with large opposing signs on currencies; for Russia those currencies were pushed around by the simultaneous, historically large Fed tightening cycle, a macro factor orthogonal to the invasion. That is why Bayesian ridge is the Russia outlier (a negative full-window gap) and why it earns the highest walk-forward overfit flag: it is the expected behaviour of an unregularised linear model in a pre-period with a strong non-event macro driver, and it is the reason the ensemble uses the *median* rather than the mean. For Hormuz, whose calmer pre-period lacks an analogous confound, the five models agree to within $\sim$ 5 pp.

6.3 Inference: the in-space placebo

Table 7 and Figure 3 report the in-space placebo. For **Hormuz**, Brent ranks *first* among all placebo units under every model and both pools, with a post/pre RMSPE ratio (6.9–10.2) above every donor’s, so the permutation p -value sits at the attainable floor (0.050 in the 19-donor pool, 0.030 in the 32-donor pool). This is the strongest signal the test can give at this sample size.

Observed vs. synthetic Brent (thin grey = the five ensemble models)




Figure 2: Observed Brent (black) vs. the synthetic counterfactual. The thick blue line is the ensemble-median synthetic; thin grey lines are the five individual models; the shaded region after T_0 (dashed red) is the estimated premium. The synthetics track Brent closely pre-event and decouple sharply afterward.

For **Russia**, no model reaches the floor (XGBoost comes closest at $p = 0.20$ in the 19-donor pool, 0.15 in the 32-donor pool); the volatile 2020–22 pre-period gives donors large pre-period RMSPE too, widening the placebo distribution. This is a known low-power property of the in-space placebo at volatile pre-periods (Abadie et al., 2010), not evidence of a small effect. Russia’s magnitude is independently documented and externally validated below.

Table 7: In-space placebo permutation p -values and Brent’s post/pre RMSPE ratio (19-donor pool). Floor = $1/(N+1)$: 0.050 (19-donor), 0.030 (32-donor).

Event	Model	p (19-donor)	p (32-donor)	Brent ratio
Russia	Convex SCM	0.50	0.50	2.75
Russia	Augmented SCM	0.60	0.47	3.51
Russia	Elastic net	0.40	0.41	4.08
Russia	XGBoost	0.20	0.15	7.75
Russia	Bayesian ridge	0.60	0.53	5.88
Hormuz	Convex SCM	0.050	0.029	6.93
Hormuz	Augmented SCM	0.050	0.029	8.45
Hormuz	Elastic net	0.050	0.029	8.53
Hormuz	XGBoost	0.050	0.029	10.19
Hormuz	Bayesian ridge	0.050	0.029	10.24

6.4 Robustness

Russia magnitude / external validation. Because Russia is a known event, we validate its magnitude against the EIA’s last two *pre-invasion* Short-Term Energy Outlook forecasts (issued before 24 February 2022), which are an independent structural counterfactual sharing no information path with the SCM. Over the identical 151-day window, the ensemble-median synthetic mean (\$88.6) sits within \$3.6 of the STEO Feb-2022 forecast (\$84.9); three of five




Figure 3: In-space placebo for Hormuz (convex SCM). Each bar is one unit’s post/pre RMSPE ratio; Brent (red) is the largest, ranking first among all 19 placebo units, so the permutation test rejects at the attainable floor.

models bracket the STEO anchor to within $\pm \$4$. Two independent methods agreeing on the counterfactual is strong evidence the SCM is not producing an implausible path.

Post-window horizon sensitivity. Figure 4 reports the gap at 1/2/3/6-month and full horizons. The two events show opposite, diagnostic profiles. Russia *declines* with horizon (31% $\rightarrow$ 22%): the signature of delayed contamination, as second-round war channels switch on through 2022 and pull the synthetic up, so the full-window figure is conservative and the least-contaminated one-month estimate is $\sim 31\%$. Hormuz *rises then eases* (46% $\rightarrow$ 59% $\rightarrow$ 56%): the chokepoint premium builds in and holds, the opposite of attenuation; the slight easing to 56% at the full horizon reflects Brent’s partial mid-June pullback, captured by extending the window to the latest data. That the same diagnostic declines for Russia and rises for Hormuz shows it is not a mechanical artefact.




Figure 4: Ensemble-median gap by post-event horizon (shaded = IQR across models). Russia attenuates with horizon (delayed-contamination signature); Hormuz rises then plateaus (the premium builds in and holds).

Leave-one-donor-out. Dropping each weighted donor in turn leaves the Hormuz estimate tight around its baseline (convex 51.2–59.3, ASCM 54.6–60.5, elastic 55.2–62.2, XGBoost 51.6–57.7%); no single donor drives the result. Russia is wider (e.g. convex 32.8–59.6%), consistent with its volatile pre-period and the simplex’s reliance on a few soft-commodity donors.

Cross-event weight transfer. Applying Russia-fitted weights to the Hormuz window (Table 8), convex SCM is the only model that transfers to a clearly *positive* Hormuz premium (25.3%) with a degraded but non-pathological pre-fit (pre-RMSPE $5\times$ independent); the augmented SCM transfers to roughly zero (−0.8%, $7\times$), and the signed-regression, tree, and prior-dominated models go strongly negative with pre-RMSPE 14–280 $\times$ independent. We therefore read the transfer as weak, convex-only corroboration that *some* positive premium survives a change of regime, not as confirmation of its level: the headline magnitude rests on the within-event ensemble. Once the window is extended through Brent’s June pullback, the regime gap between the 2020–22 and 2024–26 factor structures is evidently wider than over the shorter window.

Table 8: Cross-event weight transfer: Russia-fitted weights applied to Hormuz. Independent and transferred pre-period RMSPE in log units.

Model	Indep. gap (%)	Transf. gap (%)	Indep. RMSPE	Transf. RMSPE
Convex SCM	55.9	25.3	0.066	0.336
Augmented SCM	56.8	−0.8	0.066	0.467
Elastic net	56.6	−68.6	0.054	1.513
XGBoost	51.2	−44.1	0.053	0.928
Bayesian ridge	43.0	−100.0	0.037	18.567

Naïve counterfactual floor. Against donor-free univariate counterfactuals (Table 9), the SCM adds value, and the naïve methods’ bias *flips sign across events* purely because the pre-window trend sign differs. A linear trend extrapolates Russia’s recovery rally upward and returns an implausible *negative* full-window gap (−4%), but extrapolates Hormuz’s mild downtrend and *inflates* the gap to 71%. The trend-free random-walk null gives Russia +8.7% and Hormuz +48.9%, with the SCM ~ 13 and ~ 7 pp above respectively. The SCM, anchored to donor co-movement rather than Brent’s own momentum, is far more stable across the two regimes, which is the concrete demonstration of why a donor-based counterfactual is necessary.

Table 9: SCM vs. naïve univariate counterfactuals, full-window gap (%). RW flat/drift: random-walk carry / with drift; Lin. trend: linear pre-trend extrapolation; ARIMA rw/mr: frozen AR(1) forecasts (random-walk / mean-reverting).

Event	SCM median	RW flat	RW drift	Lin. trend	ARIMA rw	ARIMA mr
Russia	22.0	8.7	−6.9	−4.2	9.2	26.5
Hormuz	55.9	48.9	49.8	70.9	48.1	48.7

Sign of the contamination bias. Using equation (3) with a Brent-free common factor, we sign the bias per (linear) model. For **Hormuz** the full-window bias is small and *negative* (conservative) for all four linear models (−1.3 to −7.6%), so the large premium is not a contamination artefact. For **Russia** the full-window bias is large and *positive* (overestimate, +13 to +21% for three of four models), concentrated in low-reliability soft commodities whose 2022 idiosyncratic drift the convex/elastic fits leaned on; following the reliability logic of Ham and Miratrix (2024), splitting donors by pre-window factor R^2 shows the bias sits mostly in the low- R^2 (unreliable) donors.

The asymmetric conclusion is that contamination is conservative for the focal Hormuz estimate but inflates the Russia *full-window* number; Russia’s magnitude validation therefore rests on the early/peak window (where the \$95→\$130 move is recovered), not on the attenuated full-window average.

6.5 Interpretation

Our estimate places the 2026 Hormuz chokepoint premium at roughly 56 % of Brent’s no-event level over the first ~ 3.5 months. The number holds up across all five estimators and across the checks we ran: the in-space placebo reaches its floor, leave-one-out stays tight, the contamination bias is conservative, and the SCM clears the naïve baselines. The cross-event transfer is the weakest of the checks: only convex SCM transfers to a clearly positive premium, so we lean on it for the sign rather than the level. We read the estimate as a complement to, not a competitor of, the structural scenario estimates from energy institutions: those answer “what price does an assumed barrel shortfall imply,” while the SCM answers “what price gap does the observed cross-asset co-movement imply.” For the policy levers of Section 2 (reserve releases, monetary pass-through, fiscal buffers), the integrated, path-level premium is the relevant input, and the estimate here provides it with an auditable identification argument.

7 Conclusion

We estimated the Brent-price effect of the 2026 Strait of Hormuz disruption by building a counterfactual with an ensemble synthetic control, validating the design first on the known Russia 2022 event. Over the window to mid-June 2026 the ensemble-median premium is $\sim 56\%$ (IQR 51–57%), large and positive under every estimator, and it survives the in-space placebo (Brent first among all units, p at the floor), leave-one-out, the contamination-bias sign check, and the naïve-baseline comparison; the cross-event transfer gives weaker, convex-only support for a positive premium. On the calibration event the same machinery recovers a magnitude consistent with the documented \$95→\$130 move and within \$3.6/bbl of the EIA’s pre-invasion structural forecast.

The estimate’s value is as a reduced-form, data-driven complement to structural scenario analysis, with an identification argument a reviewer can audit donor by donor. Several limitations qualify it. The Hormuz post-window is short (~ 72 observations), so every inference test is under-powered and the placebo can only reach its floor; the convex synthetic under-represents Brent’s volatility, which the augmented SCM corrects; the cross-event transfer degrades once the window is extended through Brent’s June pullback; the Bayesian-ridge model is a regression proxy for BSTS, not the full state-space model; and the estimand is a reduced-form total effect, not a supply/demand/risk decomposition. Natural extensions are a longer post-window as data accrue, per-donor delayed-contamination onset dating to sharpen the post-window choice, and a full BSTS implementation for autocorrelation-aware credible intervals. None of these changes the central finding: the observed co-movement of global markets with Brent implies that the 2026 Hormuz disruption added a large, statistically distinguishable premium to crude, on the order of half its no-event price.

References

Alberto Abadie. Using synthetic controls: Feasibility, data requirements, and methodological aspects. *Journal of Economic Literature*, 59(2):391–425, 2021.

- Alberto Abadie and Javier Gardeazabal. The economic costs of conflict: A case study of the Basque country. *American Economic Review*, 93(1):113–132, 2003.
- Alberto Abadie and Jérémy L’Hour. A penalized synthetic control estimator for disaggregated data. *Journal of Business & Economic Statistics*, 39(2):492–504, 2021.
- Alberto Abadie, Alexis Diamond, and Jens Hainmueller. Synthetic control methods for comparative case studies: Estimating the effect of California’s tobacco control program. *Journal of the American Statistical Association*, 105(490):493–505, 2010.
- Alberto Abadie, Alexis Diamond, and Jens Hainmueller. Comparative politics and the synthetic control method. *American Journal of Political Science*, 59(2):495–510, 2015.
- Donald W. K. Andrews. Tests for parameter instability and structural change with unknown change point. *Econometrica*, 61(4):821–856, 1993.
- Susan Athey, Mohsen Bayati, Nikolay Doudchenko, Guido Imbens, and Khashayar Khosravi. Matrix completion methods for causal panel data models. *Journal of the American Statistical Association*, 116(536):1716–1730, 2021.
- Jushan Bai and Pierre Perron. Estimating and testing linear models with multiple structural changes. *Econometrica*, 66(1):47–78, 1998.
- Jushan Bai and Pierre Perron. Computation and analysis of multiple structural change models. *Journal of Applied Econometrics*, 18(1):1–22, 2003.
- Christiane Baumeister and James D. Hamilton. Structural interpretation of vector autoregressions with incomplete identification: Revisiting the role of oil supply and demand shocks. *American Economic Review*, 109(5):1873–1910, 2019.
- Martin Becker, Gregor Pfeifer, and Karsten Schweikert. Price effects of the Austrian fuel price fixing act: A synthetic control study. *Energy Economics*, 97:105207, 2021.
- Eli Ben-Michael, Avi Feller, and Jesse Rothstein. The augmented synthetic control method. *Journal of the American Statistical Association*, 116(536):1789–1803, 2021.
- Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B*, 57(1):289–300, 1995.
- Christoph Bergmeir and José M. Benítez. On the use of cross-validation for time series predictor evaluation. *Information Sciences*, 191:192–213, 2012.
- Kay H. Brodersen, Fabian Gallusser, Jim Koehler, Nicolas Remy, and Steven L. Scott. Inferring causal impact using Bayesian structural time-series models. *Annals of Applied Statistics*, 9(1):247–274, 2015.
- Stephen J. Brown and Jerold B. Warner. Using daily stock returns: The case of event studies. *Journal of Financial Economics*, 14(1):3–31, 1985.
- Dario Caldara and Matteo Iacoviello. Measuring geopolitical risk. *American Economic Review*, 112(4):1194–1225, 2022.
- Ivan A. Canay, Joseph P. Romano, and Azeem M. Shaikh. Randomization tests under an approximate symmetry assumption. *Econometrica*, 85(3):1013–1030, 2017.

- Gavin C. Cawley and Nicola L. C. Talbot. On over-fitting in model selection and subsequent selection bias in performance evaluation. *Journal of Machine Learning Research*, 11:2079–2107, 2010.
- Qihui Chen and Ge Yan. A mixed placebo test for synthetic control method. *Economics Letters*, 224:111004, 2023.
- Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016.
- Gregory C. Chow. Tests of equality between sets of coefficients in two linear regressions. *Econometrica*, 28(3):591–605, 1960.
- Nikolay Doudchenko and Guido W. Imbens. Balancing, regression, difference-in-differences and synthetic control methods: A synthesis. NBER Working Paper 22791, National Bureau of Economic Research, 2016.
- Federal Reserve Bank of Dallas. The impact of the 2026 Iran war on U.S. inflation. Working Paper 2609, Federal Reserve Bank of Dallas, 2026.
- Orkideh Gharehgozli. An estimation of the economic cost of recent sanctions on Iran using the synthetic control method. *Economics Letters*, 157:141–144, 2017.
- Jinyong Hahn and Ruoyao Shi. Synthetic control and inference. *Econometrics*, 5(4):52, 2017.
- Dae Woong Ham and Luke Miratrix. Benefits and costs of matching prior to a difference in differences analysis when parallel trends does not hold. *arXiv preprint arXiv:2205.08644*, 2024.
- James D. Hamilton. Causes and consequences of the oil shock of 2007–08. *Brookings Papers on Economic Activity*, pages 215–261, 2009.
- David I. Harvey, Stephen J. Leybourne, and A. M. Robert Taylor. Simple, robust, and powerful tests of the breaking trend hypothesis. *Econometric Theory*, 25(4):995–1029, 2009.
- Rob J. Hyndman and George Athanasopoulos. *Forecasting: Principles and Practice*. OTexts, 2nd edition, 2018.
- International Energy Agency. Oil market report (april and may 2026). Technical report, International Energy Agency, 2026.
- Luis Juárez, Carlos Chávez, and Mauricio Rojas. When prices go awry: Synthetic control estimates of the regional impact of the 2014 oil-price bust. *Working Paper*, 2024.
- Kiel Institute for the World Economy. The cost of closing the strait of Hormuz. Kiel Policy Brief 206, Kiel Institute for the World Economy, 2026.
- Lutz Kilian. Not all oil price shocks are alike: Disentangling demand and supply shocks in the crude oil market. *American Economic Review*, 99(3):1053–1069, 2009.
- Denis Kwiatkowski, Peter C. B. Phillips, Peter Schmidt, and Yongcheol Shin. Testing the null hypothesis of stationarity against the alternative of a unit root. *Journal of Econometrics*, 54(1-3):159–178, 1992.
- Michele Mancini, Francesco Paolo Conteduca, and Alessandro Borin. The real effects of the Russia–Ukraine war on trade flows. *World Trade Review*, 23(4):433–447, 2024.

Oxford Institute for Energy Studies. Modelling the Brent price response to a Strait of Hormuz disruption. Energy comment, Oxford Institute for Energy Studies, 2026.

Ashesh Rambachan and Jonathan Roth. A more credible approach to parallel trends. *Review of Economic Studies*, 90(5):2555–2591, 2023.

Jonathan Roth. Pretest with caution: Event-study estimates after testing for parallel trends. *American Economic Review: Insights*, 4(3):305–322, 2022.

Gábor J. Székely, Maria L. Rizzo, and Nail K. Bakirov. Measuring and testing dependence by correlation of distances. *Annals of Statistics*, 35(6):2769–2794, 2007.

U.S. Energy Information Administration. Europe brent spot price fob (series RBRTed), daily. <https://www.eia.gov/dnav/pet/hist/RBRTed.htm>, 2026.

World Bank. Commodity markets outlook. Technical report, World Bank, 2026.

Hui Zou and Trevor Hastie. Regularization and variable selection via the elastic net. *Journal of the Royal Statistical Society: Series B*, 67(2):301–320, 2005.

A Robustness grid

The headline uses the preferred specification (shared 19-donor pool, preferred pre-window, ensemble median). The full grid varies (i) the pre-window (preferred, extended, and narrow specifications); (ii) the donor pool (shared 19, permissive 26 for Russia, full 32); and (iii) the model (each of the five plus the median). For Russia the donor-pool sweep is itself a diagnostic: the gap is expected to *shrink* monotonically from the clean 19-donor pool to the contaminated 32-donor pool, because Russia-treated donors get inflated by the same shock and pull the synthetic up. If this monotonic pattern holds, the SUTVA-driven exclusion logic is empirically supported.

B Bayesian ridge and full BSTS

The fifth ensemble member implements only the regression-on-donors term $\beta'x_t$ of the Bayesian structural time-series model of Brodersen et al. (2015), with a conjugate Gaussian–Gamma prior. It omits the local-linear *trend*, the *seasonal* component, the *spike-and-slab* donor selection, and the autocorrelation-aware posterior over the full state trajectory. At our sample size, with no strong weekly seasonality in daily oil prices, the regression term carries most of the signal, so the proxy is an adequate Bayesian-flavoured diversifier; but its credible intervals are Gaussian and ignore autocorrelation, and its donor “importances” are standardised-coefficient magnitudes, not true posterior inclusion probabilities. We therefore call it a Bayesian ridge regression rather than a BSTS, and read its placebo p -values as rank statistics rather than formal probabilities (Table 7). A full BSTS implementation is the natural route to autocorrelation-aware credible intervals in a future iteration.

C Structural-break tests

Two break tests support donor screening. *Bai–Perron* (Bai and Perron, 1998, 2003) dates breaks in the donor–Brent weekly-return relationship $r_t^{\text{donor}} = \alpha_s + \beta_s r_t^{\text{Brent}} + \varepsilon_t$ by globally minimising segment SSR via dynamic programming, with the break count chosen by BIC. Only VIX shows

a break strictly inside a pre-window (2022-01-07). We treat the break tests as informative rather than automatically disqualifying: the VIX break is a loading-stability signal, it does not reproduce on the analysis window alone, and the in/out sensitivity is below 1 pp, so VIX is retained (Section 4). The relationship specification is Brent-based, so a break *at* T_0 is ambiguous; the complementary [Harvey et al. \(2009\)](#) trend-break test resolves this by testing the donor’s *own* log-level trend with a statistic whose asymptotic size is invariant to whether the series is $I(0)$ or $I(1)$: it weights an $I(0)$ -optimal and an $I(1)$ -optimal sup- t statistic by KPSS stationarity weights ([Kwiatkowski et al., 1992](#)). Empirically all 32 donors are classified $I(1)$ -like and none rejects a trend break within ± 8 weeks of either T_0 , so no retained donor was itself trend-treated by either event.